

## CPU Performance

- **System Configuration**
  - Memory : RAM 4GB
  - CPU core: 2
  - CPU Threads : 2
  - Disk :40 GB
  - m1-medium
  - CentOS7

CPU benchmarking for Integer and Floating point operations

Number of Thread	GIOPS	GFLOPS
1	10.519813	9.990418
2	19.351904	19.189053
4	19.072120	19.005313
8	19.700419	18.593352


Table 1: Normal Instructions

Number of Thread	GIOPS	GFLOPS
1	3.004218	3.817306
2	5.997491	7.627489
4	11.218634	14.224442
8	19.912195	27.375987


Table 2: AVX Instructions

### Analysis:

- The above two tables and the below graph shows the speed of the CPU in terms of Giga Integer Operations(GIOPS) and Giga Floating Point Operations (GFLOPS) for execution of normal instructions and for Advanced Vector Instructions(AVX instructions).
- The experiment is performed with varying number of threads – 1,2,4,8 and the GFLOPS and GIOPS are calculated.
- The Table 1, shows that as the number of threads is increasing for The GIOPS and GFLOPS are also increasing. For the benchmarking with normal instructions, the IOPS and FLOPS increases and then after thread 4, it is constant at 19 GIOPS.



- The Table 2, which is for AVX, we can see the similar results, but the number of GIOPS and GFLOPS are small initially and then rises gradually to 19 as the thread increases.
- The data in both the tables is plotted graphically, GIOPS for normal instructions is compared with GIOPS for AVX Instructions, and similarly for GFLOPS.
- We can see from both the graphs that as the number of thread increases the AVX instructions performs better while normal instructions run at stable speed.



- **Theoretical peak performance**  

$$= \text{CPU Frequency} * \text{Number of Cores} * \text{Threads per core} * \text{IPC}$$

$$= 2.93 * 2 * 2 * 8$$

$$= 93.76 \text{ GFLOPS}$$
- Our CPU Benchmark compared to the theoretical performance have an efficiency 20.45% for normal instructions.

#### LINPACK BENCHMARK

- The linpack benchmark uses basic vector operation and matrix operations.
- The problem size used is 20000. The average performance for this size is 72.35 GFLOPS. The best performance is 72.54 GFLOPS for the problem size of 20000.
- Compared to Theoretical performance, the linpack have an efficiency of 77.16%.
- Compared to our performance we achieved an efficiency of 26.51%, considering GFLOPS for 2 threads as we get best result for 2 threads.
- If we consider the GFLOPS of AVX instructions we achieve an efficiency of 37.75% considering 8 thread GFLOPS.

```

cc@pai-rohan:~/l_mklb_p_2018.0.006/benchmarks_2018/linux/mkl/benchmarks/linpack
Sun Oct 8 15:38:35 UTC 2017
Sample data file lininput_xeon64.

Current date/time: Sun Oct 8 15:38:35 2017

CPU frequency: 3.089 GHz
Number of CPUs: 2
Number of cores: 2
Number of threads: 2

Parameters are set to:

Number of tests: 15
Number of equations to solve (problem size) : 1000 2000 5000 10000 15000 18000 20000 22000 25000 26000 27000 30000 35000 40000 45000
Leading dimension of array : 1000 2000 5000 10000 15000 18000 20016 22008 25000 26000 27000 30000 35000 40000 45000
Number of trials to run : 4 2 2 2 2 2 2 2 2 2 1 1 1 1 1
Data alignment value (in Kbytes) : 4 4 4 4 4 4 4 4 4 4 1 1 1 1 1

Maximum memory requested that can be used=3202964416, at the size=20000

===== Timing linear equation system solver =====

Size LDA Align. Time(s) GFlops Residual Residual(norm) Check
1000 1000 4 0.022 29.7489 9.394430e-13 3.203742e-02 pass
1000 1000 4 0.015 43.8587 9.394430e-13 3.203742e-02 pass
1000 1000 4 0.014 47.5642 9.394430e-13 3.203742e-02 pass
1000 1000 4 0.012 54.2088 9.394430e-13 3.203742e-02 pass
2000 2000 4 0.103 52.0376 4.085732e-12 3.554086e-02 pass
2000 2000 4 0.086 62.1631 4.085732e-12 3.554086e-02 pass
5000 5000 4 1.219 68.3893 2.262585e-11 3.154992e-02 pass
5000 5000 4 1.195 69.7582 2.262585e-11 3.154992e-02 pass
10000 10000 4 9.121 73.1158 9.187981e-11 3.239775e-02 pass
10000 10000 4 9.410 70.8644 9.187981e-11 3.239775e-02 pass
15000 15000 4 31.079 72.4115 2.219450e-10 3.495671e-02 pass
15000 15000 4 30.741 73.2070 2.219450e-10 3.495671e-02 pass
18000 18000 4 54.569 71.2612 2.886628e-10 3.161212e-02 pass
18000 18000 4 53.421 72.7928 2.886628e-10 3.161212e-02 pass
20000 20016 4 73.910 72.1703 3.669736e-10 3.248520e-02 pass
20000 20016 4 73.526 72.5475 3.669736e-10 3.248520e-02 pass

Performance Summary (GFlops)

Size LDA Align. Average Maximal
1000 1000 4 43.8451 54.2088
2000 2000 4 57.1003 62.1631
5000 5000 4 69.0738 69.7582
10000 10000 4 71.9901 73.1158
15000 15000 4 72.8092 73.2070
18000 18000 4 72.0270 72.7928
20000 20016 4 72.3589 72.5475

Residual checks PASSED

End of tests

```

**MEMORY PERFORMANCE:**



- For performing memory benchmark, we have used strong scaling experiment, where in we have allocated 1GB of memory and the operations are performed.

**SEQUENTIAL READ+WRITE****Throughput (MB/s)**

Number of Thread	8B	8KB	8MB	80MB
1	1434.84	41433.8184	8785.41	8348.7
2	2802.42	37720.2572	10400	8128.31
4	3110.26	45547.6837	13485.6	7315.88
8	3067.17	37534.937	13653.2	5394.38

**Latency(microsec)**

Number of Thread	8B	8KB	8MB	80MB
1	0.00066	2.3017E-05	0.00011	0.00011
2	0.00017	1.2641E-05	4.6E-05	5.9E-05
4	7.7E-05	5.2345E-06	1.8E-05	3.3E-05
8	3.9E-05	3.176E-06	8.7E-06	2.2E-05

**Analysis**



- The throughput is increasing when we run our benchmark from 8B block size to 8KB block but decreases further when we increase the block.
- As we can see that the throughput is increasing and latency is decreasing as the number of thread increases from 1 to 8.
- The throughput spikes up for 8kb block.
- The optimal number of thread we get is 4 as we have best performance for 8B and 8KB block and a marginal less throughput in 8MB block.

**RANDOM READ+WRITE MEMORY****Throughput (MB/s)**

Number of Thread	8B	8KB	8MB	80MB
1	330.79	32888.2477	8664.77	6980.13
2	32.6984	19499.5932	13133.9	9254.29
4	26.0631	21173.4085	12587.8	8205.8
8	33.8029	22851.8876	15114.7	5989.49

**Latency (micro sec)**

Number of Thread	8B	8KB	8MB	80MB
1	0.00288	2.9E-05	0.00011006	0.00014
2	0.01458	2.4E-05	3.6306E-05	5.2E-05
4	0.00915	1.1E-05	1.894E-05	2.9E-05
8	0.00353	5.2E-06	7.887E-06	2E-05

**Analysis**


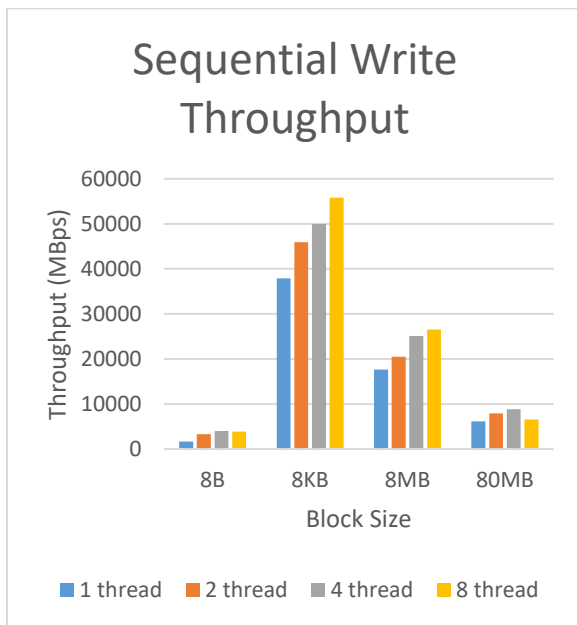
- The random read + write operation has a very less throughput and high latency for the block of 8B irrespective of the number of threads.
- The throughput spikes for 8KB block and then decreases for 8 MB and further for 80MB.
- The latency is low for 8KB, 8MB, 80MB as the block size is more and not dependent on the operation.
- The optimal number of thread is difficult to know as for 8KB block 1 thread performs better, for 8MB – number of thread = 8 performs better and for 80 MB thread 2 gives maximum throughput.

**Sequential READ+WRITE MEMORY****Throughput (MBps)**

Number of Thread	8B	8KB	8MB	80MB
1	1687.07	37906	17618.9064	6106.05
2	3329.7	45955.9	20511.5485	7881.59
4	4004.97	50002.4	25121.1586	8817.42
8	3834.49	55800.5	26516.1874	6547.21

**Latency(micro sec)**

Number of Thread	8B	8KB	8MB	80MB
1	0.00057	2.5E-05	5.4128E-05	0.00016
2	0.00014	1E-05	2.3247E-05	6.1E-05
4	6E-05	4.8E-06	9.4907E-06	2.7E-05
8	3.1E-05	2.1E-06	4.4957E-06	1.8E-05

**Analysis**



- The sequential write operation performs similar to sequential read + write operation.
- The throughput spikes up for 8KB block and the latency is low for this block.
- The number of thread when 4 and 8 gives good throughput performance
- Also as the number of threads are increasing the throughput is increasing as well.

### Random Write Throughput (MBps)

Number of Thread	8B	8KB	8MB	80MB
1	327.779	35221.1	16321.9638	7547.74
2	26.96	21750.5	18678.7085	9009.89
4	24.085	31408.1	24715.8473	9319.14
8	28.01	28925.8	24870.02	6325.59

### Latency(microsec)

Number of Thread	8B	8KB	8MB	80MB
1	0.00291	2.9E-05	5.8429E-05	0.00013
2	0.01769	2.2E-05	2.5528E-05	5.3E-05
4	0.0099	7.6E-06	9.6464E-06	2.6E-05
8	0.00426	4.1E-06	4.7933E-06	1.9E-05



### Analysis

- The random write operation results are similar to random read+write operation.
- The throughput spikes for 8KB block and then decreases for 8 MB and further for 80MB.
- The latency is low for 8KB, 8MB, 80MB as the block size is more and not dependent on the operation.
- The optimal number of thread is difficult to know as for 8KB block 1 thread performs better, for 8MB and 80MB – number of thread = 4 gives better results.

## THEORETICAL BENCHMARK

= Clock Frequency \* Number of data transfer per clock \* Memory bus interface width \* number of interfaces

= 2.93GHz \* 2 \* 64bits \* 2

= 75.08GHz

## STREAM BENCHMARK

```

cc@pa1-rohan:~
[cc@pa1-rohan ~]$ gcc stream.c
[cc@pa1-rohan ~]$ ./a.out
-----
STREAM version $Revision: 5.10 $
-----
This system uses 8 bytes per array element.
-----
Array size = 10000000 (elements), Offset = 0 (elements)
Memory per array = 76.3 MiB (= 0.1 GiB).
Total memory required = 228.9 MiB (= 0.2 GiB).
Each kernel will be executed 10 times.
The *best* time for each kernel (excluding the first iteration)
will be used to compute the reported bandwidth.
-----
Your clock granularity/precision appears to be 1 microseconds.
Each test below will take on the order of 29601 microseconds.
    (= 29601 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----
WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----
Function      Best Rate MB/s  Avg time     Min time     Max time
Copy:         5016.1    0.032914    0.031897    0.034762
Scale:        5001.9    0.032903    0.031988    0.034167
Add:          6565.7    0.037518    0.036554    0.039572
Triad:        6318.0    0.038945    0.037987    0.040959
-----
Solution Validates: avg error less than 1.000000e-13 on all three arrays
-----
[cc@pa1-rohan ~]$

```

The best rate achieved for performing memory copy function is 5016.1MB/s.


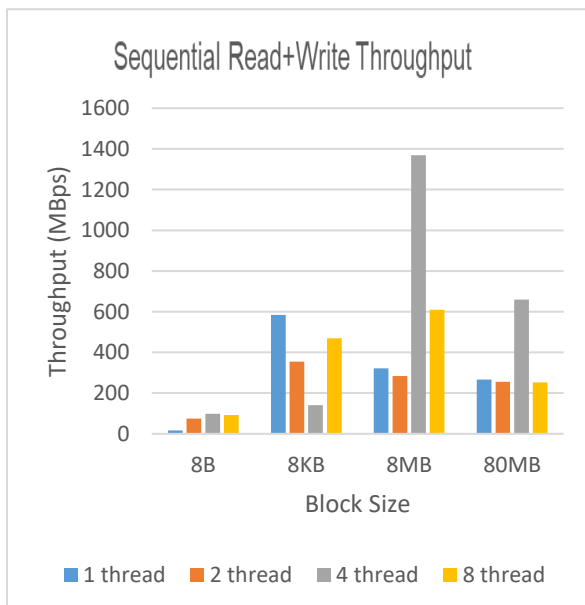


**DISK BENCHMARK****Sequential Read Write  
Throughput(MBps)**

Number of Thread	8B	8KB	8MB	80MB
1	16.241194	584.147248	321.853274	266.554358
2	75.327317	354.749373	283.295392	255.383030
4	98.338451	141.230566	1367.958705	660.344544
8	92.918141	469.685636	609.796716	251.735583

**Latency(ms)**

Number of Thread	8B	8KB	8MB	80MB
1	61.571826219559	1.711897134781	3.107005834579	3.751579999924
2	13.275396466255	2.818891525269	3.529884457588	3.915686964989
4	10.168962240219	7.080620229244	0.731016218662	1.514360964298
8	10.762161105871	2.129083633423	1.639890760183	3.972422122955

**Analysis**

- We performed disk operations read+write on the 10GB data.

- We can see that with low block size we are having less throughput and as the block size increases the throughput increases as well.
- As the number of thread increases the throughput also increases irrespective of block size.
- The optimal number of thread is 4 for sequential read+write operation.


### Random read Write


#### Throughput (MBps)

Number of Thread	8B	8KB	8MB	80MB
1	6.853093	2543.196956	3071.499082	160.652898
2	9.105492	2840.872547	4291.521993	129.973646
4	8.715221	2819.311624	4153.177163	439.406512
8	8.186703	2905.573539	4202.278348	827.351493

#### Latency(ms)

Number of Thread	8B	8KB	8MB	80MB
1	145.919523954391	0.393205881119	0.325573921204	6.224599838257
2	109.823823451996	0.352004528046	0.233017563820	7.693867444992
4	114.741778731346	0.354696512222	0.240779519081	2.275796949863
8	122.149295389652	0.344166129827	0.237966150045	1.208676129580





#### Analysis


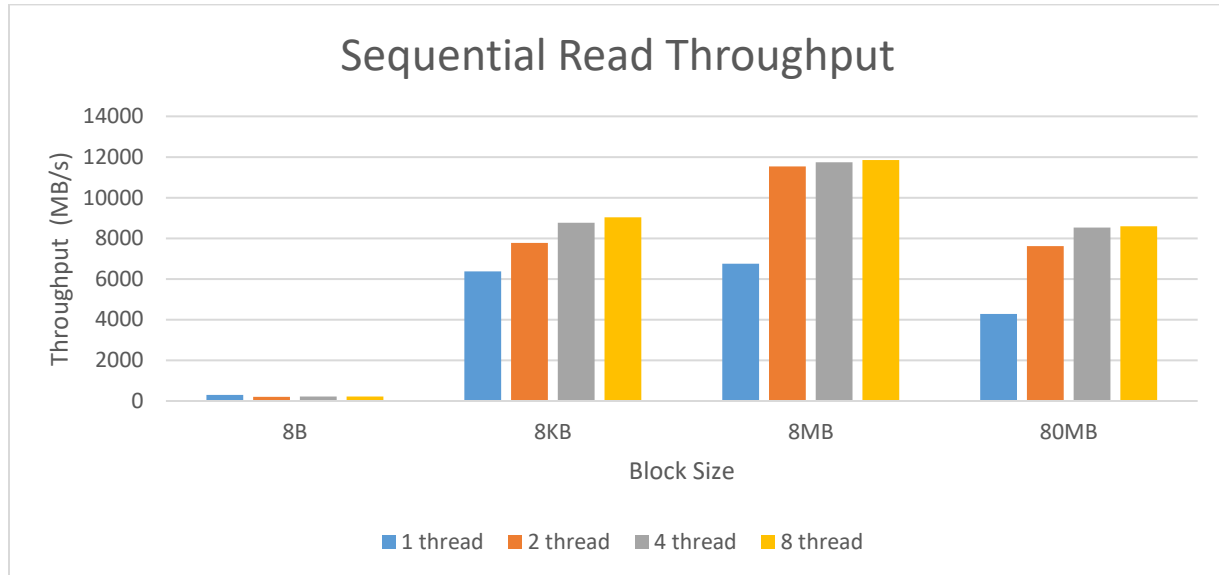
- The random read +write operation results are similar to sequential read + write operations.
- The throughput increases with increase in blocks size and also with increase in thread.
- The latency is high in 8B block access.
- The optimal number of thread is 8 for random read+write operation.

#### Sequential Read Throughput (MB/s)

Number of Thread	8B	8KB	8MB	80MB
1	306.777814	6368.853492	6752.840447	4285.849751
2	212.102434	7781.466832	11538.551898	7622.791841
4	214.969523	8766.624864	11740.644065	8523.184827
8	217.693971	9028.695081	11856.773546	8598.664782

#### Latency (ms)

Number of Thread	8B	8KB	8MB	80MB
1	3.259688138962	0.157014131546	0.148085832596	0.233325958252
2	4.714703083038	0.128510475159	0.086665987968	0.131185531616
4	4.651822209358	0.114068984985	0.085174202919	0.117327034473
8	4.593604475260	0.110757976770	0.084339976311	0.116297125816



#### Analysis

- The throughput increases and latency decreases with the increase in block size and also the with the increase in thread.
- Only for 80MB block the throughput and latency starts decreasing.

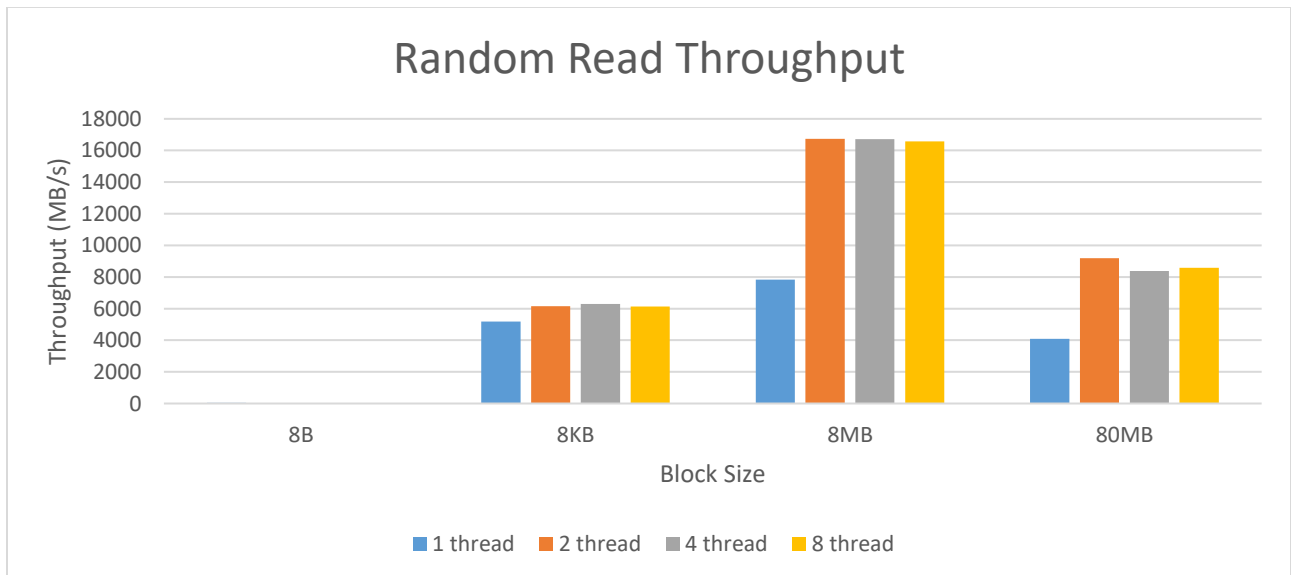
- The optimal number of thread can be 4 or 8 as the sequential read operation have almost equal throughput for all the blocks.


### Random Read Throughput (MB/s)

Number of Thread	8B	8KB	8MB	80MB
1	36.715080	5177.213518	7841.110313	4086.537909
2	30.257610	6154.373485	16726.467896	9182.700381
4	31.194412	6299.125185	16725.634144	8386.079597
8	29.357849	6133.182330	16571.931202	8590.365622

### Latency (ms)

Number of Thread	8B	8KB	8MB	80MB
1	27.236765146255	0.193154096603	0.127532958984	0.244705915451
2	33.049536466599	0.162486076355	0.059785485268	0.108900427818
4	32.057023227215	0.158752202988	0.059788465500	0.119245231152
8	34.062440991402	0.163047492504	0.060342997313	0.116409480572





#### Analysis

- The throughput increases and latency decreases with the increase in block size and also the with the increase in thread.
  - Only for 80MB block the throughput and latency starts decreasing.
  - The optimal number of thread can be 2 or 4 as the random read operation have almost equal throughput for all the blocks.
- 
- The optimal number of concurrency to obtain best performance is 4 or 8 number of threads.
  - Our chameleon instances has SSD disk.

```

iozone: Performance Test of File I/O
Version $Revision: 3.394 $
Compiled for 64 bit mode.
Build: linux

Contributors: William Norcott, Don Capps, Isom Crawford, Kirby Collins
Al Slater, Scott Rhine, Mike Wisner, Ken Goss
Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,
Ben England.

Run began: Sat Oct 7 05:14:37 2017

Auto Mode
Command line used: ./iozone -a
Output is in Kbytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 Kbytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

KB      reclen      write      rewrite      read      reread      read      write      read      record      stride
64      4      581273 1310683 2561267 4018152 277290 2662899 3791156 3541098 4272462
64      8 1484662 3541098 5389653 15972885 9318832 3203069 5283570 4564786 2892445
64      16 735831 3738358 6421025 15972885 9318832 3363612 5389653 1991276 7100397
64      32 771797 4018152 15972885 642102512902017 3541098 5860307 4564786 2923952
64      64 1599680 4018152 6421025 6421025 4564786 1879725 5735102 1879725 6421025
128      4 1508887 1852520 4135958 10779307 8036304 3895854 5847904 4596273 6114306
128      8 1622919 1852520 5603747 1420079410567140 3895854 7476717 2511022 3867787
128      16 1997245 4717434 16365173 1588107812542043 1399360 8036304 2608629 9287508
128      32 1526043 4889281 5545860 1588107812842051 5603747 8036304 5784891 9129573
128      64 1727352 2166499 15881078 640613814200794 4934216 3359523 6114306 9977956
128      128 1997245 2202044 6114306 15881078 6114306 3895854 7476717 3657616
256      4 886667 2849590 4281170 4350555 2588541 2639446 6936061 5455847
256      8 1312954 4929815 13454450 1416439511091721 5455847 9114515 6936061
256      16 1610276 5455847 8815202 6107548 5217259 2392442 9114515 7314033
256      32 856386 2170027 14953435 149534351454450 639872010245071 7571923
256      64 1003679 2539563 16072608 1607260814164395 693606110651598 7791708
256      128 1829808 2533571 14164395 5971678 8534922 7120034 4197489 5320671

KB      reclen      write      rewrite      read      reread      read      write      read      record      stride
64      4      581273 1310683 2561267 4018152 277290 2662899 3791156 3541098 4272462
64      8 1484662 3541098 5389653 15972885 9318832 3203069 5283570 4564786 2892445
64      16 735831 3738358 6421025 15972885 9318832 3363612 5389653 1991276 7100397
64      32 771797 4018152 15972885 642102512902017 3541098 5860307 4564786 2923952
64      64 1599680 4018152 6421025 6421025 4564786 1879725 5735102 1879725 6421025
128      4 1508887 1852520 4135958 10779307 8036304 3895854 5847904 4596273 6114306
128      8 1622919 1852520 5603747 1420079410567140 3895854 7476717 2511022 3867787
128      16 1997245 4717434 16365173 1588107812542043 1399360 8036304 2608629 9287508
128      32 1526043 4889281 5545860 1588107812842051 5603747 8036304 5784891 9129573
128      64 1727352 2166499 15881078 640613814200794 4934216 3359523 6114306 9977956
128      128 1997245 2202044 6114306 15881078 6114306 3895854 7476717 3657616
256      4 886667 2849590 4281170 4350555 2588541 2639446 6936061 5455847
256      8 1312954 4929815 13454450 1416439511091721 5455847 9114515 6936061
256      16 1610276 5455847 8815202 6107548 5217259 2392442 9114515 7314033
256      32 856386 2170027 14953435 149534351454450 639872010245071 7571923
256      64 1003679 2539563 16072608 1607260814164395 693606110651598 7791708
256      128 1829808 2533571 14164395 5971678 8534922 7120034 4197489 5320671



65536 1024 2144069 3490610 5973554 6249241 6306447 3754352 6284675 8738265 6744509 2982687 2944096 6042892 6352209
65536 2048 2147738 3530647 6227861 6398639 6463030 3764274 6250520 8625296 6924782 2665927 2769918 6080050 6318624
65536 4096 2445776 3776635 7683775 7788938 7581624 4067035 7855044 8635865 12656799 3274039 3084917 7619664 7658727
65536 8192 2490810 3777517 8103758 8209524 7967982 4071311 7743082 8489177 12208219 3503467 3320992 7822408 7865608
65536 16384 2354781 3514532 6151897 6252226 6175115 3841716 6116171 4696615 6972737 3837425 3472136 6050740 6338147
131072 64 2471514 3738517 8135032 8180670 8165481 3628612 7929802 10693806 8250280 3727238 3686647 8067464 8068411
131072 128 2481532 3684596 7733366 7865136 7839676 4073206 7652417 10478963 7770092 3559999 3558248 7934151 7965997
131072 256 2472403 758248 7155036 7317899 7346258 4006125 7162027 9435793 7318776 3479757 3474699 7438802 7438399
131072 512 2491113 7430233 7492841 7565366 7532366 3968605 7639550 8511099 7552440 3502526 3474391 7419426 7432064
131072 1024 2495551 3803701 7835989 7596165 7528549 4048311 7508704 8531704 7877421 3285757 3264625 7831747 
```

**NETWORK BENCHMARK****TCP - Loopback**

Number of Thread	Throughput(Mbps)	Latency(ms)
1	1598.401598	500.499999
2	1196.54189	751.492123
4	830.57617	1252.8129
8	539.56229	2255.5287

**TCP – Two Node**

Number of Thread	Throughput(Mbps)	Latency(ms)
1	1252.936	638.5
2	1196.5793	751.597
4	834.259	1250.25
8	542.125	2251.25

**Analysis**

- For TCP we are sending 100 MB file from client to server and back to the client.
- As we can see that as the number of threads are increasing the throughput is decreasing for both loopback and two nodes.





**UDP - Loop Back Interface**

Number of Thread	Throughput(Mbps)	Latency(ms)
1	7944.38927507448	1006.9999
2	16235.413495687468	985.5
4	28070.175438596492	1140.0
8	32837.35248845562	1949.0

**UDP – Two Nodes (192.168.0.195) and (192.168.0.216)**

Number of Thread	Throughput(Mbps)	Latency(ms)
1	736.750011511719	10858.5
2	1016.8091258619046	15735.5
4	1061.7296239154598	30139.5
8	1297.9243350672791	49309.5

**Analysis**

- We performed UDP for 1GB data transfer from client to server and back to client with block size of 64KB.
- For loopback, as the number of threads are increasing the network speed is also increasing.
- While for two nodes the network speed is almost stable over all the threads.
- Same thing we can observe for latency. As the throughput is increasing in loopback, the latency is almost stable across all the threads.
- And as the throughput is stable in two node, the latency increases across all the threads.

## IPERF BENCHMARK

## TCP Loopback

```
Client connecting to 127.0.0.1, TCP port 5001
TCP window size: 2.50 MByte (default)
-----
[  3] local 127.0.0.1 port 36644 connected with 127.0.0.1 port 5001
[  5] local 127.0.0.1 port 5001 connected with 127.0.0.1 port 36644
[ ID] Interval           Transfer             Bandwidth
[  3]  0.0- 5.0 sec      27045 MBytes      5409 MBytes/sec
[  3]  5.0-10.0 sec     27377 MBytes      5475 MBytes/sec
[  3] 10.0-15.0 sec     26415 MBytes      5283 MBytes/sec
[  3] 15.0-20.0 sec     28356 MBytes      5671 MBytes/sec
[  3] 20.0-25.0 sec     23863 MBytes      4773 MBytes/sec
[  3]  0.0-25.0 sec    133056 MBytes     5322 MBytes/sec
[  3] MSS size 65468 bytes (MTU 65508 bytes, unknown interface)
[  5]  0.0-25.0 sec     130 GBytes      44.6 Gbits/sec
```

## UDP Loopback

```
Client connecting to 127.0.0.1, UDP port 5001
Sending 1470 byte datagrams, IPG target: 1.18 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[  3] local 127.0.0.1 port 40482 connected with 127.0.0.1 port 5001
[ ID] Interval           Transfer             Bandwidth
[  3]  0.0-10.0 sec     10.1 GBytes      8.71 Gbits/sec
[  3] Sent 7408438 datagrams
[  3] WARNING: did not receive ack of last datagram after 10 tries.
```